

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf

from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from tensorflow.keras import layers

import os
import shutil

import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: DATASET_PATH = 'Automating_Port_Operations_dataset'

# Courtesy of SimpliLearn for providing os commands regarding train_test_split on d
def copy_files(img_paths, img_classes, target):
    for img_path, img_class in zip(img_paths, img_classes):
        dest_dir = os.path.join(target, img_class)
        os.makedirs(dest_dir, exist_ok=True)
        shutil.copy(img_path, dest_dir)

def train_test_dir(test_size, rand, data_dir=DATASET_PATH, train_dir=DATASET_PATH+"
os.makedirs(train_dir, exist_ok=True)
os.makedirs(test_dir, exist_ok=True)

path_list = []
classes = []

for folder in os.listdir(data_dir) :
    folder_path = os.path.join(data_dir, folder)
    if os.path.isdir(folder_path) :
        for file in os.listdir(folder_path) :
            if file.endswith('.jpg') :
                path_list.append(os.path.join(folder_path, file))
                classes.append(folder)

train_paths, test_paths, train_classes, test_classes = train_test_split(path_li

copy_files(train_paths, train_classes, train_dir)
copy_files(test_paths, test_classes, test_dir)
```

## CNN

```
In [ ]: # Batch size, image size
BAT_SIZE, IMG_SIZE = 32, 180

train_test_dir(0.2, 43)

train = tf.keras.preprocessing.image_dataset_from_directory(DATASET_PATH + "/train"
```

```

label_mode = 'categorical'
validation_split = 0.2,
shuffle = True,
seed = 43,
subset = 'training',
image_size = (IMG_SIZE,
batch_size = BAT_SIZE)
val = tf.keras.preprocessing.image_dataset_from_directory(DATASET_PATH + "/train",
label_mode = 'categorical'
validation_split = 0.2,
shuffle = True,
seed = 43,
subset = 'validation',
image_size = (IMG_SIZE,
batch_size = BAT_SIZE)

test = train = tf.keras.preprocessing.image_dataset_from_directory(DATASET_PATH + "
label_mode = 'categorical'
shuffle = True,
seed = 43,
image_size = (IMG_SIZE,
batch_size = BAT_SIZE)

class_names = train.class_names

AUTOTUNE = tf.data.AUTOTUNE

train = train.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val = val.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 1083 files belonging to 9 classes.  
Using 867 files for training.  
Using 867 files for training.  
Found 1083 files belonging to 9 classes.  
Using 216 files for validation.  
Found 503 files belonging to 9 classes.

```

In [ ]: model = tf.keras.models.Sequential([
layers.Rescaling(1./255, input_shape = (IMG_SIZ
layers.Conv2D(32, 3, padding='same', activation
layers.MaxPooling2D(),
layers.Conv2D(32, 3, padding='same', activation
layers.MaxPooling2D(),
layers.GlobalAveragePooling2D(),
layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dense(128, activation='relu'),
layers.Dense(9, activation='softmax'))])

model.compile(optimizer='adam',
loss='categorical_crossentropy',
metrics=['accuracy', 'precision', 'recall'])


# Uncomment for model layout
# print(model.summary())

hist = model.fit(train,


```

```
        validation_data=val,  
        epochs = 20)  
  
acc = hist.history['accuracy']  
v_acc = hist.history['val_accuracy']  
loss = hist.history['loss']  
v_loss = hist.history['val_loss']
```


Epoch 1/20

**16/16**  **5s** 176ms/step - accuracy: 0.2523 - loss: 2.1046 - precision: 0.1765 - recall: 0.0018 - val\_accuracy: 0.3380 - val\_loss: 1.9057 - val\_precision: 0.4478 - val\_recall: 0.1389


Epoch 2/20

**16/16**  **2s** 135ms/step - accuracy: 0.3425 - loss: 1.8286 - precision: 0.4745 - recall: 0.0947 - val\_accuracy: 0.3380 - val\_loss: 1.8466 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00


Epoch 3/20

**16/16**  **2s** 124ms/step - accuracy: 0.3589 - loss: 1.7776 - precision: 0.4941 - recall: 0.0078 - val\_accuracy: 0.3380 - val\_loss: 1.8323 - val\_precision: 0.4000 - val\_recall: 0.0093


Epoch 4/20

**16/16**  **2s** 121ms/step - accuracy: 0.3101 - loss: 1.8079 - precision: 0.0882 - recall: 7.0167e-04 - val\_accuracy: 0.3426 - val\_loss: 1.8138 - val\_precision: 0.6000 - val\_recall: 0.0417


Epoch 5/20

**16/16**  **2s** 124ms/step - accuracy: 0.3774 - loss: 1.7365 - precision: 0.7921 - recall: 0.0597 - val\_accuracy: 0.3472 - val\_loss: 1.8038 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00


Epoch 6/20

**16/16**  **2s** 124ms/step - accuracy: 0.3778 - loss: 1.7417 - precision: 0.0980 - recall: 3.5878e-04 - val\_accuracy: 0.3611 - val\_loss: 1.8028 - val\_precision: 0.2222 - val\_recall: 0.0093


Epoch 7/20

**16/16**  **2s** 123ms/step - accuracy: 0.3820 - loss: 1.6404 - precision: 0.5324 - recall: 0.0315 - val\_accuracy: 0.3426 - val\_loss: 1.7746 - val\_precision: 0.5000 - val\_recall: 0.0231


Epoch 8/20

**16/16**  **2s** 117ms/step - accuracy: 0.4286 - loss: 1.6337 - precision: 0.6077 - recall: 0.0542 - val\_accuracy: 0.3889 - val\_loss: 1.7531 - val\_precision: 0.5556 - val\_recall: 0.0694


Epoch 9/20

**16/16**  **2s** 112ms/step - accuracy: 0.4278 - loss: 1.6613 - precision: 0.5290 - recall: 0.0746 - val\_accuracy: 0.3935 - val\_loss: 1.7336 - val\_precision: 0.5000 - val\_recall: 0.0231


Epoch 10/20

**16/16**  **2s** 115ms/step - accuracy: 0.4136 - loss: 1.6515 - precision: 0.5236 - recall: 0.0432 - val\_accuracy: 0.3750 - val\_loss: 1.7533 - val\_precision: 0.5577 - val\_recall: 0.1343


Epoch 11/20

**16/16**  **2s** 114ms/step - accuracy: 0.4126 - loss: 1.6833 - precision: 0.6054 - recall: 0.1091 - val\_accuracy: 0.3981 - val\_loss: 1.7369 - val\_precision: 0.5714 - val\_recall: 0.1111


Epoch 12/20

**16/16**  **2s** 114ms/step - accuracy: 0.4137 - loss: 1.7069 - precision: 0.5357 - recall: 0.0770 - val\_accuracy: 0.4028 - val\_loss: 1.7213 - val\_precision: 0.5636 - val\_recall: 0.1435


Epoch 13/20

**16/16**  **2s** 113ms/step - accuracy: 0.4460 - loss: 1.6598 - precision: 0.5668 - recall: 0.1502 - val\_accuracy: 0.4028 - val\_loss: 1.7118 - val\_precision: 0.6538 - val\_recall: 0.0787


Epoch 14/20

**16/16**  **2s** 112ms/step - accuracy: 0.3859 - loss: 1.6415 - precision: 0.5880 - recall: 0.1146 - val\_accuracy: 0.4028 - val\_loss: 1.6959 - val\_precision: 0.6471 - val\_recall: 0.1019


Epoch 15/20

**16/16**  **2s** 116ms/step - accuracy: 0.4365 - loss: 1.5789 - precision: 0.5766 - recall: 0.1098 - val\_accuracy: 0.3981 - val\_loss: 1.6844 - val\_precision: 0.6111 - val\_recall: 0.1019


Epoch 16/20

**16/16**  **2s** 129ms/step - accuracy: 0.4395 - loss: 1.6296 - precision: 0.5523 - recall: 0.0733 - val\_accuracy: 0.4074 - val\_loss: 1.6819 - val\_precision: 0.6279 - val\_recall: 0.1250


Epoch 17/20

**16/16**  **2s** 129ms/step - accuracy: 0.4368 - loss: 1.5475 - precision: 0.6522 - recall: 0.1615 - val\_accuracy: 0.4074 - val\_loss: 1.6679 - val\_precision: 0.6765 - val\_recall: 0.1065


Epoch 18/20

**16/16**  **2s** 126ms/step - accuracy: 0.4220 - loss: 1.6160 - precision: 0.7559 - recall: 0.1277 - val\_accuracy: 0.4074 - val\_loss: 1.6717 - val\_precision: 0.6094 - val\_recall: 0.1806

Epoch 19/20

**16/16**  **2s** 126ms/step - accuracy: 0.4012 - loss: 1.5983 - precision: 0.5835 - recall: 0.1418 - val\_accuracy: 0.4074 - val\_loss: 1.6589 - val\_precision: 0.6897 - val\_recall: 0.0926

Epoch 20/20

**16/16**  **2s** 126ms/step - accuracy: 0.4210 - loss: 1.5599 - precision: 0.6939 - recall: 0.1686 - val\_accuracy: 0.3750 - val\_loss: 1.7194 - val\_precision: 0.5893 - val\_recall: 0.1528


```
In [ ]: plt.figure(figsize=(12,6))

plt.subplot(1, 2, 1)
plt.plot(range(20), acc, label='Train Accuracy')
plt.plot(range(20), v_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Train vs Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(range(20), loss, label='Train Loss')
plt.plot(range(20), v_loss, label='Validation Loss')
plt.legend(loc='lower left')
plt.title('Train vs Validation Loss')

plt.suptitle('CNN Model')
plt.show()
```

## CNN Model



```
In [ ]: predicts = model.predict(test)
        predict = np.argmax(predicts, axis=1)
        true = np.argmax(tf.concat([y for x, y in test], axis=0), axis=1)

        # print(f"\nAccuracy score: {accuracy_score(true, predict)}")
        # print(f"Mean Squared Error: {mean_squared_error(true, predict)}\n")

        model.evaluate(test)
```

16/16 ————— 1s 54ms/step

16/16 ————— 1s 56ms/step - accuracy: 0.4173 - loss: 1.5982 - precision: 0.5923 - recall: 0.1523

```
Out[ ]: [1.6276891231536865,
         0.3956262469291687,
         0.5813953280448914,
         0.14910537004470825]
```

```
In [ ]: import seaborn as sns

        sns.heatmap(confusion_matrix(true, predict), annot=True, fmt='d',
                     xticklabels=class_names, yticklabels=class_names)
        plt.title('CNN Heatmap')
        plt.show()

        print(classification_report(true, predict))
```



```

shuffle = True,
seed = 1,
subset = 'training',
image_size = (IMG_SIZE,
batch_size = BAT_SIZE)
val = tf.keras.preprocessing.image_dataset_from_directory(DATASET_PATH + "/train",
label_mode = 'categorical',
validation_split = 0.3,
shuffle = True,
seed = 1,
subset = 'validation',
image_size = (IMG_SIZE,
batch_size = BAT_SIZE)
test = train = tf.keras.preprocessing.image_dataset_from_directory(DATASET_PATH + "
label_mode = 'categorical',
shuffle = True,
seed = 1,
image_size = (IMG_SIZE,
batch_size = BAT_SIZE)

AUTOTUNE = tf.data.AUTOTUNE

train = train.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val = val.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 1083 files belonging to 9 classes.  
Using 759 files for training.  
Using 759 files for training.  
Found 1083 files belonging to 9 classes.  
Using 324 files for validation.  
Found 503 files belonging to 9 classes.

```

In [ ]: mn_model = tf.keras.applications.MobileNetV2(input_shape = (IMG_SIZE, IMG_SIZE, 3),
include_top=False)

mn_model.trainable = False

# Uncomment for MobileNet model summary
# print(mn_model.summary())

# inputs
i = tf.keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))

# hidden
x = layers.Rescaling(1./255)(i)
x = mn_model(x, training=False)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.2)(x)
x = layers.Flatten()(x)
x = layers.Dense(256, activation='relu')(x)
x = layers.BatchNormalization()(x)
x = layers.Dropout(0.1)(x)
x = layers.Dense(128, activation='relu')(x)
x = layers.BatchNormalization()(x)
x = layers.Dropout(0.1)(x)

#outputs

```



```
o = layers.Dense(9, activation='softmax')(x)

model = tf.keras.Model(i, o)

# Uncomment for final model summary
# print(model.summary())

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy', 'precision', 'recall'])

callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3)

hist = model.fit(train,
                 validation_data = val,
                 epochs = 50,
                 callbacks=[callback])

acc = hist.history['accuracy']
v_acc = hist.history['val_accuracy']
loss = hist.history['loss']
v_loss = hist.history['val_loss']
```

Epoch 1/50  
**16/16** ————— **11s** 359ms/step - accuracy: 0.4118 - loss: 1.7864 - precision: 0.5920 - recall: 0.2750 - val\_accuracy: 0.6358 - val\_loss: 1.1212 - val\_precision: 0.8962 - val\_recall: 0.5062  
Epoch 2/50  
**16/16** ————— **4s** 274ms/step - accuracy: 0.8185 - loss: 0.6302 - precision: 0.9207 - recall: 0.7229 - val\_accuracy: 0.7284 - val\_loss: 0.8148 - val\_precision: 0.8734 - val\_recall: 0.6389  
Epoch 3/50  
**16/16** ————— **4s** 261ms/step - accuracy: 0.9018 - loss: 0.3420 - precision: 0.9547 - recall: 0.8400 - val\_accuracy: 0.8025 - val\_loss: 0.6178 - val\_precision: 0.8931 - val\_recall: 0.7222  
Epoch 4/50  
**16/16** ————— **4s** 262ms/step - accuracy: 0.9527 - loss: 0.1844 - precision: 0.9732 - recall: 0.9256 - val\_accuracy: 0.8580 - val\_loss: 0.5140 - val\_precision: 0.9184 - val\_recall: 0.7994  
Epoch 5/50  
**16/16** ————— **4s** 257ms/step - accuracy: 0.9781 - loss: 0.1265 - precision: 0.9836 - recall: 0.9462 - val\_accuracy: 0.8735 - val\_loss: 0.4695 - val\_precision: 0.9215 - val\_recall: 0.8333  
Epoch 6/50  
**16/16** ————— **4s** 260ms/step - accuracy: 0.9907 - loss: 0.0873 - precision: 0.9988 - recall: 0.9840 - val\_accuracy: 0.8704 - val\_loss: 0.4740 - val\_precision: 0.9130 - val\_recall: 0.8426  
Epoch 7/50  
**16/16** ————— **4s** 262ms/step - accuracy: 1.0000 - loss: 0.0592 - precision: 1.0000 - recall: 0.9949 - val\_accuracy: 0.8827 - val\_loss: 0.4372 - val\_precision: 0.9200 - val\_recall: 0.8519  
Epoch 8/50  
**16/16** ————— **4s** 258ms/step - accuracy: 0.9891 - loss: 0.0556 - precision: 0.9955 - recall: 0.9863 - val\_accuracy: 0.8889 - val\_loss: 0.4280 - val\_precision: 0.9123 - val\_recall: 0.8673  
Epoch 9/50  
**16/16** ————— **4s** 259ms/step - accuracy: 1.0000 - loss: 0.0423 - precision: 1.0000 - recall: 0.9984 - val\_accuracy: 0.8920 - val\_loss: 0.4574 - val\_precision: 0.9150 - val\_recall: 0.8642  
Epoch 10/50  
**16/16** ————— **4s** 256ms/step - accuracy: 0.9977 - loss: 0.0326 - precision: 0.9977 - recall: 0.9977 - val\_accuracy: 0.8580 - val\_loss: 0.4975 - val\_precision: 0.8980 - val\_recall: 0.8426  
Epoch 11/50  
**16/16** ————— **4s** 264ms/step - accuracy: 0.9872 - loss: 0.0395 - precision: 0.9919 - recall: 0.9852 - val\_accuracy: 0.8457 - val\_loss: 0.5593 - val\_precision: 0.8766 - val\_recall: 0.8333

```
In [ ]: model.evaluate(test)

num_iters = len(hist.history['loss'])

plt.figure(figsize=(12,6))


plt.subplot(1, 2, 1)
plt.plot(range(num_iters), acc, label='Train Accuracy')
plt.plot(range(num_iters), v_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Train vs Validation Accuracy')
```

```
plt.subplot(1, 2, 2)
plt.plot(range(num_iters), loss, label='Train Loss')
plt.plot(range(num_iters), v_loss, label='Validation Loss')
plt.legend(loc='lower left')
plt.title('Train vs Validation Loss')

plt.suptitle('MobileNet Model')
plt.show()
```

16/16 ————— 3s 170ms/step - accuracy: 0.9989 - loss: 0.0238 - precision: 0.9989 - recall: 0.9964

MobileNet Model



## Closing Statement

MobileNet's model can achieve more than twice the accuracy & a much less loss score than base CNN's model in less epochs. This is due to:

1. MobileNet having a more elaborate model - not only in having extra layers for dropout, but in the transfer model itself.
2. MobileNet is already a prefixed model that has been transferred over, meaning it is capable of being more accurate and precise than the crudely assembled CNN model.

However, base CNN's model has less of a divergence in the train-validation loss score compared to MobileNet. For example, there is a 0.1 difference in loss in the base CNN, compared to MobileNet's 0.6 difference. This is perhaps a drawback to MobileNet's more scrutinized model, with the extra layers causing a higher variation in score.